

## Relational Model

- **Domain:** Set of atomic values
- **Relation:** Set of tuples
- **Superkey:** Subset of attributes that uniquely identifies one tuple
- **Key:** Minimal superkey
- **Candidate Keys:** Set of all keys
- **Primary Key:** Chosen candidate key; cannot be NULL
- **Foreign Key:** Refers to a primary key in another relation; must appear as a primary key of another relation or have NULL for an attribute

## Relational Algebra

$\sigma_{[c]}$ (R): Select

- Select rows that satisfy condition c
- Principle of Acceptance

$\pi_{[I]}$ (R): Project

- Select columns listed in I

$\rho_{[B_1 \leftarrow A_1, B_2 \leftarrow A_2]}(R)$ : Rename

- Order does not matter
- No two attributes can be renamed to the same name
- No attributes can be renamed more than once in a single operation

### Set Operations

- Relations must be union compatible:
  1. Have same number of attributes
  2. Attributes have same or compatible domains

### Cross Product

- Set of attributes must be disjoint
- $|R \times S| = |R| \times |S|$
- If either R or S is empty, result is an empty relation

### Join

- **Inner Joins**
  - $\bowtie_{\theta}$  : choose if tuples satisfy the condition
  - $\bowtie_{=}$  : choose if tuples satisfy the condition; condition only uses =
  - $\bowtie$  : choose if all common attributes between two tuples are equal. Becomes cross product if no common attributes
  - Common attributes (i.e. columns) can appear twice in output relation, unless natural join is used (then only appear once)
- **Outer Joins**
  - $\bowtie_{\bowtie} R \bowtie S$ , plus dangling tuples in R. Dangling tuples have values NULL for attributes from S
  - $\bowtie_{\bowtie} R \bowtie S$ , plus dangling tuples in S. Dangling tuples have values NULL for attributes from R
  - $\bowtie_{\bowtie} R \bowtie S$ , plus dangling tuples in R and S

### Equivalence

- Strong equivalence: Both queries produce error or both queries always produce the same results
- Weak equivalence: Both queries always produce the same results if neither queries produce an error

Joins	Select and Project
$R \times S \neq S \times R$ $R \bowtie S \neq S \bowtie R$ $(R \times S) \times T \equiv R \times (S \times T)$ $(R \bowtie S) \bowtie T \equiv R \bowtie (S \bowtie T)$ $(R \bowtie_{\theta} S) \bowtie_{\theta} T \neq R \bowtie_{\theta} (S \bowtie_{\theta} T)$ (weak)	$\sigma_{c_1}(\sigma_{c_2}(R)) \equiv \sigma_{c_2}(\sigma_{c_1}(R))$ $\sigma_{c_1}(\sigma_{c_2}(R)) \equiv \sigma_{c_1 \wedge c_2}(R)$ $\pi_{\ell_1}(\pi_{\ell_2}(R)) \neq \pi_{\ell_2}(\pi_{\ell_1}(R))$ (weak, unless $\ell_1 \subseteq \ell_2$ ) $\pi_{\ell}(\sigma_{\theta}(R)) \neq \sigma_{\theta}(\pi_{\ell}(R))$ (weak) $\sigma_{\theta}(R \times S) \neq \sigma_{\theta}(R) \times S$ (weak)
Set operations	
$R \cup S \equiv S \cup R$ $R \cap S \equiv S \cap R$ $(R \cup S) \cup T \equiv R \cup (S \cup T)$ $(R \cap S) \cap T \equiv R \cap (S \cap T)$  $R - S \neq S - R$ $\sigma_c(R - S) \equiv \sigma_c(R) - S \neq \sigma_c(R) - \sigma_c(S)$ (unless R and S have some common attribute names used in c) $\pi_{\ell}(R \cup S) \equiv \pi_{\ell}(R) \cup \pi_{\ell}(S)$	

## SQL DDL

### CREATE

```
CREATE TABLE <table_name> (  
    <attr1> <type> [<col_constraint>],  
    ...  
    [<table_constraints>]  
);
```

### INSERT

```
INSERT INTO <table_name> [(attr1, attr2,...)]  
VALUES (val1, val2,...), ...;
```

- Either all are inserted, or none inserted

### DELETE

```
DELETE FROM <table_name>  
[WHERE <conditions>];
```

- Delete all tuples that match given condition; if no condition, all tuples are deleted
- Principle of **ACCEPTANCE**

### Integrity Constraints

- Principle of **REJECTION**

Common constraints		
	Column	Table
Primary Key	PRIMARY KEY	PRIMARY KEY (attr1, attr2)
Unique	UNIQUE	UNIQUE (attr1, attr2)
Foreign Key	REFERENCES R(attr1, attr2)	FOREIGN KEY (attr1, attr2) REFERENCES R(attr3, attr4)
CHECK	CHECK(attr <op>)	CHECK(attr <op>)

- Foreign Keys **ON UPDATE/ON DELETE** specifies behaviour of referencing table when data in referenced table is deleted/updated
  - **NO ACTION** : reject update/delete if violates constraint
  - **RESTRICT** : NO ACTION, but not deferrable
  - **CASCADE** : propagate delete/update to referencing tuples
  - **SET DEFAULT** : Set FK in referencing tuples to default values; default values must be PK in the referenced table
  - **SET NULL** : Set FK in referencing tuples to NULL; affected columns must not have NOT NULL constraint
- UNIQUE constraints check individual attributes using <>; 2 tuples are unique if either one contains NULL

### ALTER

- Useful for circular references ( $FK_R \rightarrow PK_S, FK_S \rightarrow PK_R$ )

```
ALTER TABLE <table_name>
```

```
[ALTER/ADD/DROP] [COLUMN/CONSTRAINT] <name>  
<changes>
```

### DROP TABLE

```
DROP TABLE [IF EXISTS]
```

```
<table_name> [, <table_name2>, ...]
```

```
[CASCADE]
```

### DEFERRABLE CONSTRAINTS

- **NOT DEFERRABLE** : (default). Constraints checked at end of SQL statement and aborts if violated
- **DEFERRABLE INITIALLY DEFERRED** : Constraints checked on COMMIT, can be temporarily violated in transaction
- **DEFERRABLE INITIALLY DEFERRED** : Constraints initially not deferrable, but can be set to deferrable later with **SET CONSTRAINTS <name> DEFERRED**
- Transaction: **BEGIN ... COMMIT;**

## SQL DQL

```
SELECT [DISTINCT] <attrs>
```

```
FROM <relations>
```

```
WHERE <conditions>
```

- Aliasing: column AS alias
- Operations
  - Maths: +, \*, /, ^, %, etc.
  - String: || (concatenate), LOWER(s), UPPER(s), etc.
  - Date Time: +, NOW(), etc.
- Principle of **ACCEPTANCE**
- = and <> can be safely used if you do not want NULL values, else use IS NULL
- Regex **LIKE <regex>**
  - `_`: Any single character
  - `%`: Any sequence of 0 or more characters

### UNION/INTERSECT/EXCEPT

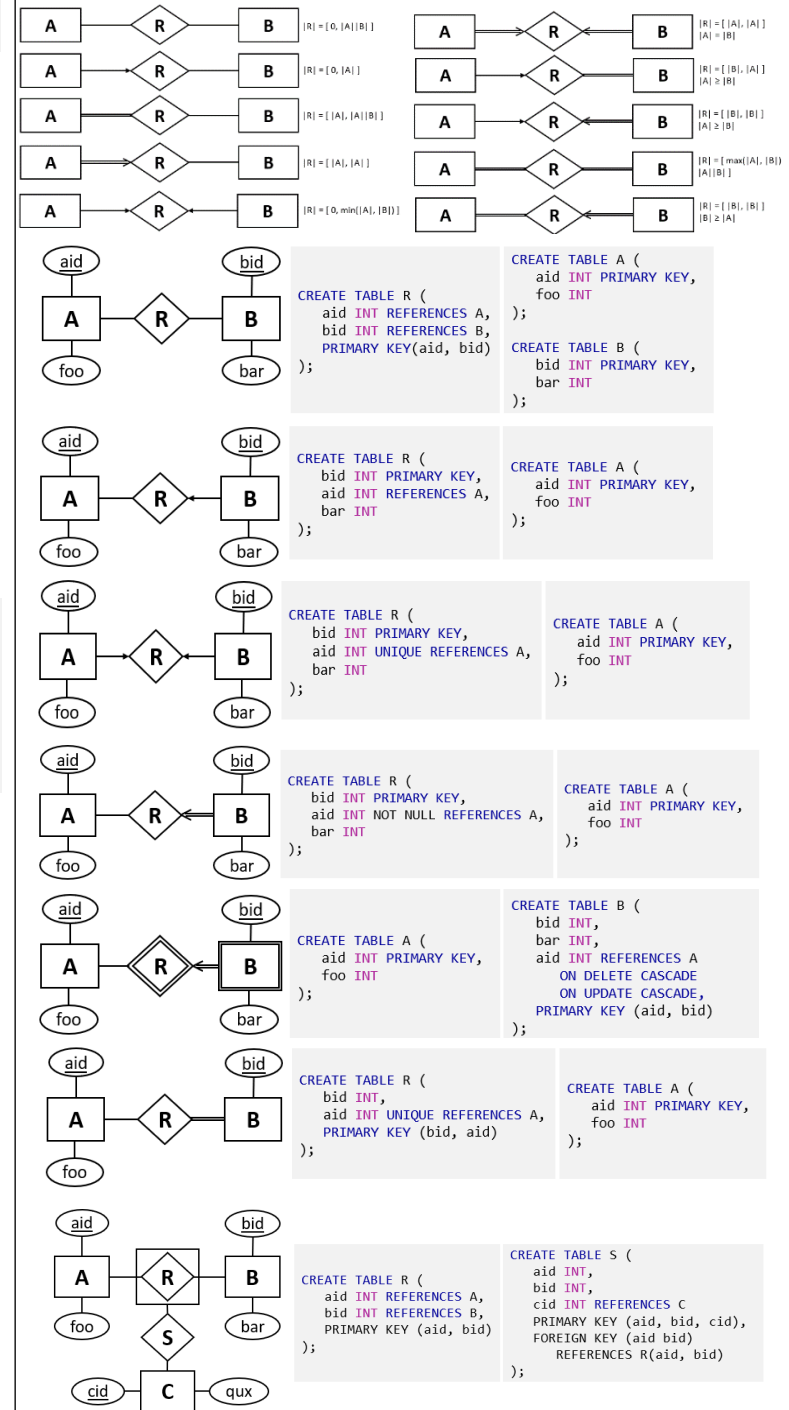
- Must be union compatible
- **UNION ALL** : # dups = #dup in R + #dup in S
- **INTERSECT ALL** : # dups = min{#dup in R, #dup in S}
- **EXCEPT ALL** : # dups = #dup in R - #dup in S

### JOIN

- Cross product **FROM R1 [AS][Alias], R2, R3,...**
  - Set of attributes need not be disjoint
- **JOIN R JOIN S ON <cond>**
- **NATURAL JOIN R NATURAL JOIN S**
  - Becomes cross product if no common attributes
- **OUTER JOIN R LEFT/RIGHT/FULL [OUTER] JOIN S ON <cond>**

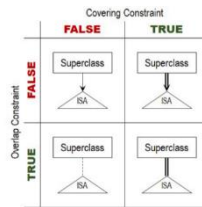
### ORDER

- Default: ASC
- **ORDER BY <attr1> DESC/ASC [, <attr2> DESC/ASC]**
- Stable sort from rightmost to left



## ISA Hierarchies

- Inherited primary key of child references the primary key of the direct parent
- Specify ON DELETE/ON UPDATE on child foreign keys



## PL/pgSQL

### Functions

CREATE OR REPLACE

FUNCTION <fn\_name> ([IN/OUT/INOUT] <param> <type> ...)

RETURNS <return\_type> AS \$\$

DECLARE --plpgsql

...

BEGIN --plpgsql

...

END;

\$\$ LANGUAGE <sql/plpgsql>;

- **IN/OUT/INOUT** :
  - **IN**: (default) Parameter is an input. A constant; cannot be reassigned
  - **OUT**: Parameter is return value. An uninitialised variable; must be assigned a value later
  - **INOUT**: Parameter is both an input and return value. An initialised variable; should be but not value later
- Return types
  - **<table\_name>**: Returns one existing tuple from the table
  - **SETOF <table\_name>**: Returns one or more existing tuples from the table
  - **RECORD**: Returns one new tuple from the table containing the **OUT/INOUT** attributes specified in the parameters list
  - **SETOF RECORD**: Returns one or more new tuples containing the **OUT/INOUT** attributes specified in the parameters list
  - **TABLE(<attr> <type>, ...)**: Returns a new table with the specified schema. Parameters list should not contain **OUT/INOUT**
    - o plpgsql: Function should call one or more **RETURN NEXT** to populate the table

### Procedures

CREATE OR REPLACE

PROCEDURE <procedure\_name> (<param> <type> ...)

AS \$\$

...

\$\$ LANGUAGE <sql/plpgsql>;

CALL procedure\_name(params...)

### IF ELSE & Control Flow

IF <cond> THEN

...

ELSE

...

END IF;

LOOP

EXIT WHEN <cond>

...

END LOOP

## Cursor

DECLARE

curs CURSOR FOR <table\_name>;

r RECORD

...

BEGIN

OPEN curs;

LOOP

FETCH curs INTO r; --get current tuple

EXIT WHEN NOT FOUND; --exit when end of table

...

RETURN NEXT; --insert tuple into table

END LOOP;

CLOSE curs;

...

## Triggers

CREATE TRIGGER <trigger\_name>

AFTER/BEFORE INSERT/UPDATE/DELETE OR [...] ON <table>

[DEFERRABLE INITIALLY DEFERRED/IMMEDIATE]

FOR EACH ROW/STATEMENT

[WHEN <cond>]

EXECUTE FUNCTION <fn\_name>();

CREATE OR REPLACE FUNCTION <fn\_name>

RETURNS TRIGGER ...

- Special variables
  - **NEW**: INSERT/UPDATE: the new tuple; DELETE: NULL
  - **OLD**: DELETE/UPDATE: the old tuple; INSERT: NULL
  - **TG\_OP**: INSERT/UPDATE/DELETE
  - **TG\_TABLE\_NAME**: Table associated with the trigger
- Only AFTER and FOR EACH ROW triggers can be deferred

### Return Types

#### BEFORE

- INSERT: Null → nothing inserted; Non-null *t* → insert *t*
- UPDATE: Null → not updated; Non-null *t* → updated to *t*
- DELETE: Null → not deleted; Non-null *t* → delete (not nec. *t*)

#### AFTER

- Does not matter; just return NULL for convenience

#### INSTEAD OF

- NULL → ugnore rest of the operation on current row; Non-null → proceed as normal
- INSTEAD OF is only defined on VIEWS and ROW-LEVEL

### RAISE

- **RAISE NOTICE '...'**: Prints warning message, but does not prevent the operation
- **RAISE EXCEPTION '...'**: Prints warning message and prevents the operation

### WHEN

- No SELECT in WHEN()
- No OLD in WHEN() for INSERT
- No NEW in WHEN() for DELETE
- No WHEN for INSTEAD OF

## Order of Execution

- BEFORE statement > BEFORE row > INSTEAD OF > AFTER row > AFTER statement
- Within each category triggers are activated in alphabetical order of their names (A → Z)
- If a BEFORE row-level trigger returns NULL, then subsequent triggers on the same row are omitted

