

Stable Matching

- **Definition:** A matching is stable if no unmatched man and woman both prefer each other to their current partners

- **Gale-Shapely Algorithm:**

```
Initialize each person to be free.
while (some man is free and hasn't proposed to every woman)
    Choose such a man m
    w = 1st woman on m's list to whom m has not yet proposed
    if (w is free)
        assign m and w to be engaged
    else if (w prefers m to her fiancé m')
        assign m and w to be engaged, and m' to be free
    else
        w rejects m
}
```

- Lemma: Men propose to women in decreasing order of preference
- Lemma: Women stay engaged after the first time they got engaged
- Lemma: Women's partners keep getting better
- Lemma: Upon termination, each man is engaged to a unique woman
- Lemma: Upon termination, the matching between men and women is stable
- The algorithm returns male-optimal stable matching

Asymptotic Analysis

- **Upper:** $T(n) = O(f(n))$
If there \exists constants $c > 0$, $n_0 \geq 0$ s.t. $\forall n \geq n_0$, $T(n) \leq c \cdot f(n)$
- **Lower:** $T(n) = \Omega(f(n))$
If there \exists constants $c > 0$, $n_0 \geq 0$ s.t. $\forall n \geq n_0$, $T(n) \geq c \cdot f(n)$
- **Tight:** $T(n) = \Theta(f(n))$
If $T(n) = O(f(n))$ and $T(n) = \Omega(f(n))$

Transitivity

- $f = O/\Omega/\Theta(g)$ and $g = O/\Omega/\Theta(h) \rightarrow f = O/\Omega/\Theta(h)$

Additivity

- $f = O/\Omega/\Theta(h)$ and $g = O/\Omega/\Theta(h) \rightarrow f + g = O/\Omega/\Theta(h)$

Rule of Thumbs

$$n^n \geq n! \geq c^{kn} \geq c^n \geq n^k \geq n \log(n) \geq n \geq \sqrt{n} \geq \log(n) \geq \log(\log(n))$$
$$O(\log(n!)) = O(n \log(n))$$

Mathematical Properties

- $a^{\log b} = b^{\log a}$
- $1 + \frac{1}{2} + \frac{1}{3} + \frac{1}{4} + \frac{1}{5} + \dots + \frac{1}{n} = O(\log(n))$
- GP sum = $\frac{a(r^n - 1)}{r - 1}$, $|r| > 1$ or $\frac{a}{1 - r}$, $|r| < 1$
- AP sum = $\frac{n}{2}(2a + (n - 1)d) = \frac{n}{2}(a + l)$

Common Recurrence Relations

- $T(n) = 2T\left(\frac{n}{2}\right) + n^2 = O(n^2)$
- $T(n) = 2^k T(\sqrt[k]{n}) + c = O(\log^k(n))$
- $T(n) = 2^k T\left(n^{\frac{1}{m}}\right) + c = O\left(\log^{\frac{k}{m}}(n)\right)$

The Master Theorem

The **Master Method** depends on the following Theorem:

Theorem: Let $a \geq 1$ and $b > 1$ be constants, let $f(n)$ be a function, and let $T(n)$ be defined on the nonnegative integers by the recurrence:

$$T(n) = aT\left(\frac{n}{b}\right) + f(n)$$

Then $T(n)$ can be bounded asymptotically as follows.

1. If $f(n) = O(n^{\log_b a - \epsilon})$ for some constant $\epsilon > 0$, then $T(n) = \Theta(n^{\log_b a})$.
2. If $f(n) = \Theta(n^{\log_b a})$, then $T(n) = \Theta(n^{\log_b a} \lg n)$.
3. If $f(n) = \Omega(n^{\log_b a + \epsilon})$ for some constant $\epsilon > 0$, and if $a \cdot f\left(\frac{n}{b}\right) \leq cf(n)$ for some constant $c < 1$ and all sufficiently large n , then $T(n) = \Theta(f(n))$.

regularity condition

Graphs

	Adjacency List	Adjacency Matrix
Check if (u,v) exists	$\Theta(\deg(u)) = O(V)$	$\Theta(1)$
Enumerate all E	$\Theta(V + E) = O(V^2)$	$\Theta(V^2)$
Use cases	Sparse graphs	Dense graphs

Definitions and theorems

- **Trees:**
 - Def: An undirected graph is a tree if it is connected and acyclic
 - Theorem: In any undirected graph G with V nodes, any two statements below imply the third:
 1. G is connected
 2. G is acyclic
 3. G has $V - 1$ edges
- **Simple path:** A path where all its vertices are distinct
- **Connected Graph (Undirected):** There is a path between any vertices u and v
- **Strongly Connected Graph (Directed):** There is a path from u to v and from v to u for any pair of vertices u and v

Graph Traversals

- **BFS:**
 - Use cases
 1. Finding shortest paths
 2. Connectivity
 3. Testing bipartitenessA graph is not bipartite if it contains an odd-lengthed cycle
Use extra 'Color' array and assign color whenever a node is added to L(i+1). After BFS, check for edges for which both ends have the same color
 - 4. Testing strong connectivity in directed graphs
BFS G from S. Then, BFS G^{ev} from S. G is strongly connected if both BFS visited all vertices in G

- Properties:
 1. Edges in graph G but not in its BFS Tree T all either connect nodes in the same layer in T or connect nodes in adjacent layers
- **DFS:**
 - Use cases
 1. Finding the set of all connected components
 - Properties:
 1. For a given recursive DFS(u) call, all vertices marked "explored" between the invocation and end of this call are descendants of u in the DFS tree T
 2. If a graph G contains an edge (u, v) that is not in its DFS tree T, then one of u or v is an ancestor or of the other (i.e. diff levels)

BFS(s):

```
visited[s] ← true
visited[v] ← false for all other v
L(0) ← list containing only s
i ← 0 // layer
T ← ∅
while L[i] is not empty do
    L(i + 1) ← empty list
    for each node u ∈ L(i)
        Consider each edge (u, v)
        if visited[v] = false then
            set visited[v] = true
            add edge (u, v) to the tree T
            add v to L(i + 1)
    endif
endfor
increment i by one
endwhile
```

DFS(s):

```
visited[s] ← false for all v
S(0) ← stack containing only s
parent[] ← empty list
T ← ∅
while S is not empty do
    pop u from S
    if visited[u] = false then
        set visited[u] = true
        add (u, parent[u]) to T
        for each edge (u, v)
            push v to S
            set parent[v] to u
        endfor
    endif
endwhile
```

DAGs

- **Properties:**
 - In every DAG, there is a node v with no incoming edges
 - A graph G is a DAG if and only if it has a topological ordering
- **Finding a topological ordering:**
 - Kahn's Algorithm [$O(V + E)$]
 1. Initialise set S that contains all nodes with no incoming edges
 2. Initialise set W to count number of incoming edges for each node
 3. Repeat until S is empty:
 - 3.1. Pick any node u from S
 - 3.2. Add u to the topological order
 - 3.3. For each (u, v_i), decrement W[v_i]
 - 3.4. If W[v_i] becomes 0, add v_i to S

Greedy

Proving Techniques

Exchange argument: Show that at each step, you can exchange S's current choice with G's current choice without hurting S's quality.

Example: Interval scheduling

1. Let $G = i_1, i_2, \dots, i_k$ and $S = j_1, j_2, \dots, j_m$ for an input L
2. Let P(m) be the proposition that if S returns m number of intervals, then G also returns m number of intervals

- Base case: $P(1)$. The optimal solution has only 1 interval. Trivially, G can pick any interval, hence $P(1)$ is true
- Inductive hypothesis: $P(m)$ is true
- $f(i_1) \leq f(j_1)$ since G always chooses the request with the earliest finish time.
- Therefore, $S^* = i_1, j_2, \dots, j_{m+1}$ is also an optimal solution (explain)

- $S^{**} = j_2, j_3, \dots, j_{m+1}$ must be optimal for $L \setminus \{i_1\}$ for S^* to be optimal for L . S^{**} outputs m intervals
- By construction, G outputs i_2, \dots, i_k for $L \setminus \{i_1\}$. By the inductive hypothesis, G must output m schedules. Hence, $m = k - 1$
- Hence, $k = m + 1$. Therefore, $P(m+1)$ is true.

Structural bound: every possible solution must adhere to some min/max and show that G produces min/max

“Greedy stays ahead”: Show that at each step, G is always as good as S . Show that “Greedy stays ahead” implies optimality

Interval Scheduling

- Rule: Schedule the request with the earliest finish time

IntervalScheduling(R):

```

A ← []
visited[s] ← true
while R is not empty do
    choose  $r_i$  in  $R$  with earlier finish time
    add  $r_i$  to A
    delete  $r_j$  in  $R$  that are incompatible with  $r_i$ 
endwhile
return A

```

Pf. (Greedy Stays Ahead)

- Let i_1, \dots, i_k be the set of requests from G and j_1, \dots, j_m be the set of requests from S
- It suffices to show that if $f(i_r) \leq f(j_r)$ for all $r \leq k$, then $k \geq m$
- Proof by induction on r
 - Base case: $r = 1$. G will choose i_1 which is the request with earliest finish time
 - Inductive hypothesis: Suppose $f(i_{r-1}) \leq f(j_{r-1})$
 - $f(i_{r-1}) \leq f(j_{r-1})$, so $f(i_{r-1}) \leq s(j_r)$
 - Hence, j_r must be available for selection by G for the r^{th} request
 - Hence, $f(i_r) \leq f(j_r)$

Interval Partitioning

- Rule: Consider the resources in the order of their start time

IntervalPartitioning (R):

```

d ← 0
sort intervals in R in ascending order of start time
for j = 1 to n
    if interval j is compatible with some resource k
        schedule interval j to resource k
    else
        allocate new resource d + 1
        schedule interval j to resource d + 1
    d ← d + 1

```

Pf. (Structural Bound)

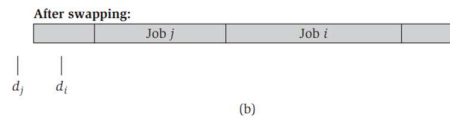
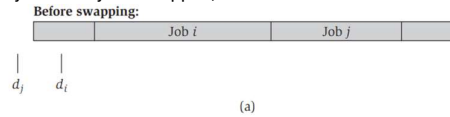
- Observation: In any instance of interval partitioning, the number of resources needed is at least the depth of the set of intervals

Minimising Lateness

- Rule: Schedule requests with the earliest deadlines first

Pf. (Exchange Argument)

- Observation: An inversion is when job j is scheduled after job i when $d(j) < d(i)$. If S has an inversion, then there must be a pair of jobs i and j such that j is scheduled immediately after i and has $d(j) < d(i)$.
- Suppose S has at least one inversion. Let job i be scheduled immediately after job j even though $d(j) < d(i)$
 - If jobs i and j are swapped, S will have one less inversion



- Denote L' as the max lateness after swap and L as the max lateness before swap. Denote $t(m)$ as the time taken to complete a job m
- $L' = \max\{t(j) - d(j), t(j) + t(i) - d(i)\}$
 $L = \max\{t(i) - d(i), t(i) + t(j) - d(j)\} = t(i) + t(j) - d(j)$, since $d(j) < d(i)$
- $t(j) - d(j) < t(i) + t(j) - d(j)$ and $t(j) + t(i) - d(i) < t(i) + t(j) - d(j)$. Therefore, L' must be smaller than L .
- We've shown that the lateness of S does not increase after the swap
- This shows that an optimal schedule with no inversions exists.
- All schedules with no inversions have the same maximum lateness (to proof). Hence, the schedule obtained by the greedy solution is optimal.

Divide and Conquer

Counting Inversions

- Mergesort, but count the number of inversions during the merge step
- Key property: if $L[i] > R[j]$, then $L[i:] > R[j:]$

Merge (L, R):

```

count ← 0
i, j ← 0
M ← empty array
while i < len(L) and j < len(R) do
    if L[i] < R[j] then
        add L[i] to back of M and increment i
    else
        add R[j] to back of M and increment j
    increment count by len(L) - i + 1
end while

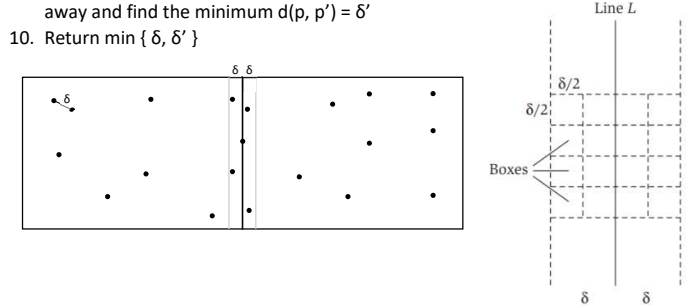
```

end while

... regular merge steps

Finding the closest pair of points

- Find closest pair in left half L and right half R and those in the boundary split
- Keep a sorted list P_x of points where points are sorted by x -axis. Keep a sorted list P_y of points where points are sorted by y -axis.
 - Split the points into two halves Q and R by their x -axis.
 - Recurse on both Q and R and find the minimum distance δ between two pairs of points between the two
 - Choose a point q in Q that has highest x value x^* and draw a vertical line L through it. L essentially divides the points in Q and R
 - (**Pf 1) Narrow down to only points with x -values inside the $[x^* - \delta, x^* + \delta]$ boundary.
 - Construct S_y containing only the points above, sorted by y -axis (can be done in $O(n)$ using P_y)
 - Divide the $[x^* - \delta, x^* + \delta]$ boundary into many $\delta/2 \times \delta/2$ boxes
 - (**Pf 2) If there exists any 2 points p and p' such that $d(p, p') < \delta$, then p and p' must be at most **15** positions away in S_y
 - Compute each $d(p, p')$ in S_y such that p and p' are within 15 positions away and find the minimum $d(p, p') = \delta'$
 - Return $\min\{\delta, \delta'\}$



- Pf 1:** If $\exists q = (q_x, q_y) \in Q$ and $r = (r_x, r_y) \in R$ s.t. $d(q, r) < \delta$, then q and r lies within distance δ of L
 $x^* - q_x \leq r_x - q_x \leq d(q, r) \leq \delta$ and $r_x - x^* \leq r_x - q_x \leq d(q, r) \leq \delta$
 \rightarrow Hence q and r have x -coordinate within δ of L
- Pf 2:** The max distance within a box is the length of the diagonal, which is $\sqrt{2 \left(\frac{\delta}{2}\right)^2} = \sqrt{\frac{\delta^2}{2}} = \frac{\delta}{\sqrt{2}} \leq \delta$. Hence, no two points can be in the same box
 Suppose there exists points s and s' in S_y such that $d(s, s') < \delta$ and that they are 16 positions apart. Assume WLOG that $s_y < s'_y$. Then, s and s' must be separated by at least 3 rows of boxes which must have a distance of at least $\frac{3}{2}\delta \geq \delta$ – a contradiction

Proving Techniques

Proof by (strong) induction on the input size n

- Define the proposition
- Show how the base case is fulfilled by the algorithm
- Suppose $P(k)$ is true for all $k < n$
- By inductive hypothesis, $P(n/c)$ must be true
- Show how combining the subproblems causes $P(n)$ to be true

Dynamic Programming

Knapsack Problem

$$DP(S, W) = \begin{cases} 0, & S = \emptyset, W \leq 0 \\ \max\{w_i + DP(S \setminus \{n_i\}, W - w_i), DP(S \setminus \{n_i\}, W)\}, & S \neq \emptyset, W > 0 \end{cases}$$

Knapsack(S_n, W):

```

M ← (n + 1) x (w + 1) array
M[0][W] ← 0 for all w
for i from 1 to n
    for w from 0 to W
        if w_i < w then
            M[i][w] ← max{ w_i + M[i - 1][w - w_i], M[i - 1][w] }
        else
            M[i][w] ← M[i - 1][w]
return M[n][W]
```

Network Flow

Definitions

- s-t cut:** partition (A, B) of V such that $s \in A$ and $t \in B$
- cap(A, B):** capacity of an s-t cut (A, B) = $\sum_{e \text{ out of } A} c(e)$
- s-t flow** must satisfy the following constraints:
 - $0 \leq f(e) \leq c(e), \forall e \in E$ (capacity)
 - $f^{in}(v) = f^{out}(v), \forall v \in V \setminus \{s, t\}$ (conservation)
- Flow value:** $v(f) = f^{out}(s)$

Max-flow Min-cut Theorem (Ford-Fulkerson Algorithm)

- Max flow value from s to t = minimum capacity of any cut
- Flow value lemma:** $f^{out}(A) - f^{in}(A) = v(f)$
 - The net flow sent across any cut in G = flow amount leaving s
- Weak duality: $v(f) \leq \text{cap}(A, B)$ for any cut (A, B)
- Certificate of Optimality: Let f be any flow and (A, B) be any cut. If $v(f) = \text{cap}(A, B)$, then f is a max flow and (A, B) is a min cut
- with the smallest capacity will be the bottleneck edge b
- For every edge e along P, if e is a forward edge, we can increase $f(e)$ by $c(b)$. If e is a backward edge, we decrease $f(e)$ by $c(b)$

FordFulkerson(V, E):

```

f(e) ← 0  ∀ e ∈ E // initialise all flows to 0
while ∃ s-t path in G do
    Find simple s-t path P in G (BFS/DFS)
    f ← Augment(f, P, G)
    Update Gf
return f
```

Augment(f, P, G):

```

b ← c(e) of bottleneck edge e along P
for each e in P
    if e is a forward edge then
        increment f(e) in G by b
    else
        decrement f(e) in G by b
return f + b
```

- Residual Graph** G_f of G can be constructed using the following rules:

- Vertices in G_f and G are the same
- For each edge e of G, if $f(e) < c(e)$, then add e to G_f but with (residual) capacity = $c(e) - f(e)$
- For each edge e of G, if $f(e) > 0$, then add e to G_f but reverse the direction

- Augmenting Path** is a simple path P in G_f from s to t.

Proof (of Ford-Fulkerson)

- Show TFAE
 - There exists a cut (A, B) s.t. $v(f) = \text{cap}(A, B)$
 - Flow f is a max flow
 - There is no augmenting path relative to f
- 1 ⇒ 2: Corollary to weak duality lemma
- 2 ⇒ 3: Proof by contrapositive
 - If there exists an augmenting path, then we can improve f by sending flow along that path (assuming capacity is a non-negative integer)
- 3 ⇒ 1:
 - Let f be a flow with no augmenting paths
 - Let A be the set of vertices reachable from s in G_f and let B = V - A
 - By defn of A, $s \in A$ and $t \notin A$ ($t \in B$)
 - Show that $f^{out}(A) = \sum_{e \text{ out of } A} c(e)$
 - There are no directed edges from A to some node u in B in G_f , otherwise vertex u would have been in set A
 - Hence, if there is an edge e from A to B in G (G now, not G_f), $f(e) = c(e)$ for e to disappear in G_f .
 - Show that $f^{in}(A) = 0$
 - Suppose $f^{in}(A) > 0$. Then, there exists an edge e(v, u) from B to A such that $v \in B$ and $u \in A$ (in G)
 - Then, there must have been a reverse edge e'(u, v) in G_f from A to B. If that's the case, v must have been reachable from s ⇒ Contradiction
 - By Flow value lemma, $v(f) = f^{out}(A) - f^{in}(A)$
 - Hence, $v(f) = \sum_{e \text{ out of } A} c(e) = \text{cap}(A, B)$ as shown above

Run time Analysis

- f(e) and c(e) are integers throughout the algorithm
- Given G(V, E), |V| = n, |E| = m, G_f will have at most 2m edges.
- Each iteration runs in $O(m+n) = O(m)$ time (dfs) and increments the flow value by at least 1.
- There is at most $v(f) \leq C$ iterations, $C = \sum_{e \text{ out of } s} c(e)$
- Hence, Ford-Fulkerson runs in $O(mC)$ time

Obtaining a min-cut

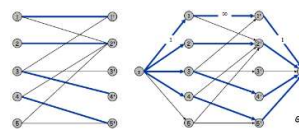
- Let A = { vertices in G_f reachable from s } and B = V - A. Then, cut (A, B) is a min-cut

Bipartite Matching

- Bipartite matching = set of edges s.t. no two edges in the set share the same endpoint. Maximum matching = largest of such set

Max flow implementation

- Add source s and join s to each node in L with edge of capacity = 1
- Add sink t and join each node in R to t with edge of capacity = 1
- Edges between L and R have infinite capacity
- Max flow in this graph G' = size of maximum matching in G



Proof (show $k \leq f$ and $k \geq f$)

Show $k \leq f$:

- Given bipartite graph G with max matching of k.
- Consider a flow f that sends 1 unit along each of the k paths
- f is a flow and has cardinality k

Show $k \geq f$:

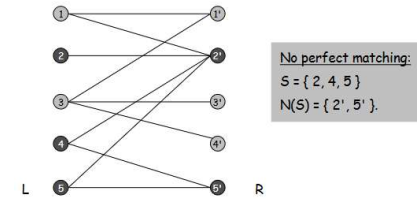
- Let f be a max flow in G' with value = k
- By integrality theorem, flow across any edge in G' is either 0 or 1
- Each node in L has at most one incoming edge from s, hence its incoming flow is at most 1 and outgoing flow (to R) is at most 1
- Each node in R has at most one outgoing edge to t, hence its outgoing flow is at most 1 and its incoming flow (from L) is at most 1
- Hence, there must be a bipartite matching between L and R with cardinality k

Perfect Matching

- Perfect matching = bipartite matching which covers all the nodes
- Condition: $|L| = |R|$

Marriage (Hall's) Theorem

- A bipartite graph $G = (L \cup R)$ has a perfect matching iff $|I'(S)| \geq |S|$ for all subsets S of L, where $I'(S)$ = set of all nodes adjacent to those in S



Proof of Marriage Theorem

⇒: G has perfect matching ⇒ $|I'(S)| \geq |S|$

Each node in S needs to be mapped to a different node in $I'(S)$ for all subsets S of L

⇐: $|I'(S)| \geq |S| \Rightarrow G$ has perfect matching. Proof by contrapositive

- Suppose G does not have a perfect matching
- Create G' from G using the same method in bipartite matching and let (A, B) be a min cut in G'
- Max flow in G' = size of max matching in G. Since size of max matching < |L| (since G has no perfect matching), $\text{cap}(A, B) < |L|$
- Let $L_A = L \cap A$, $L_B = L \cap B$, $R_A = R \cap A$
- Show that $\text{cap}(A, B) = |L_B| + |R_A|$
 - The outgoing edges from cut A must have flow = 1 because $\text{cap}(A, B) < |L|$. Hence, no edge in G (from L to R, with infinite capacity) will be part of the set of edges across the cut
 - The remaining edges are those from s to L or from R to t.
 - Edges from s to L across the cut are in L_B . Edges from R to t across the cut are from R_A
- As established in 5.1, all neighbours of nodes in L_A must be within A, so $|I'(L_A)| \leq |R_A|$
- $|I'(L_A)| \leq |R_A| = \text{cap}(A, B) - |L_B| < |L| - |L_B| = |L_A|$
- Hence, $|I'(L_A)| < |L_A|$. Choose $S = L_A$

Intractability

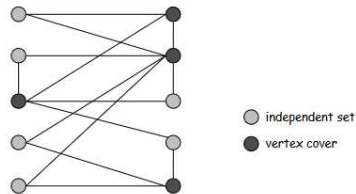
Polynomial-Time Reduction

- X polynomial reduces to Y if arbitrary instanced of X can be solved using polynomial number of calls to oracle that solved Y and polynomial number of standard computational steps (i.e. $X \leq_p Y$)
- If X reduces to Y, then Y is at least as hard as X because if Y can be solved in polynomial time, then X can be solved in polynomial time
- Intractability:** If $X \leq_p Y$ and X cannot be solved in polynomial time, then Y cannot be solved in polynomial time
- Equivalence:** If $X \leq_p Y$ and $Y \leq_p X$, then $X \equiv_p Y$

Reduction by simple equivalence

e.g. Independent set \equiv_p Vertex cover

- Independent set** = Given graph G, is there a subset of vertices S such that $|S| \geq k$ and there is no edge between each node in S
- Vertex cover** = Given graph G, is there a subset of vertices S such that $|S| \leq k$ and for each edge, at least one of its endpoints is in S



Proof (S is an independent set iff $V - S$ is a vertex cover)

\Rightarrow :

- Let S be an independent set
- Let (u, v) be an arbitrary edge
- Since S is an independent set, $u \notin S$ or $v \notin S \rightarrow u \in V - S$ or $v \in V - S$
- Hence, $V - S$ covers (u, v)

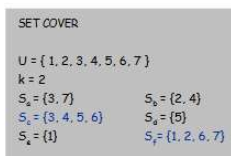
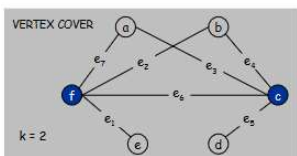
\Leftarrow :

- Let $V - S$ be a vertex cover
- Consider two nodes $u \in S$ and $v \in S$ (so $u \notin V - S$ and $v \notin V - S$)
- If u and v are joined by an edge (u, v) , then either one of u or v must be in $V - S$, which contradicts (2).
- Hence, no two nodes in S are joined by an edge $\rightarrow S$ is an independent set
- Therefore, \exists Independent set $\geq k$ iff \exists Vertex cover $< k$

Reduction from special case to general case

e.g. Vertex cover \leq_p Set cover

- Set cover:** Given a set U of elements, a collection S_1, S_2, \dots, S_m of subsets of U, is there a collection of $\leq k$ of these sets whose union is equal to U



Proof (Vertex cover reduces \leq_p Set cover)

- Construct specific set cover instance from graph G:
 - Define set U to be set of all edges E in G and each subset of edges $S_v = \{e \in E: e \text{ is incident to } v\}$
- There exists a set cover of size $\leq k$ iff there exists a vertex cover of size $\leq k$

Reduction via "gadgets"

e.g. 3-SAT \leq_p Independent set

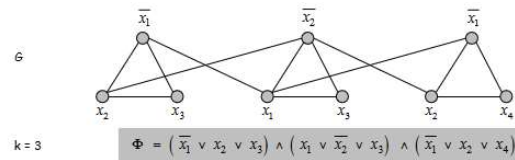
Definitions:

- Literal:** Boolean variable or its negation (x, \bar{x})
- Clause:** Disjunction of literals $(x_1 \vee \bar{x}_2 \vee x_3)$
- Conjunction normal form (CNF, Φ):** Conjunction of clauses $(x_1 \vee \bar{x}_2 \vee x_3) \wedge (x_4 \vee x_5) \wedge (x_6 \vee \bar{x}_7)$
- 3-SAT:** Given Φ where each clause contains exactly 3 literals, is there a satisfying truth assignment

Proof (3-SAT \leq_p Independent set)

Show G has an independent set S of $|S| = k = |\Phi|$ iff Φ of 3-SAT is satisfiable

- Construct specific graph G
 - G contains 3 vertices for each clause, one for each literal
 - Connect the 3 literals in a clause in a triangle via an edge
 - Connect literal to each of its negations via an edge
- \Rightarrow :
 - Let $S \subseteq G$ be an independent set of size k
 - S must contain at least one vertex from each triangle because $k = |\Phi|$
 - S must contain at most one vertex from each triangle, otherwise there would be an edge between two vertices
 - Hence, S must contain exactly one vertex from each triangle
 - Furthermore, if $u, v \in S$, then u and v cannot be negations of each other otherwise there will be an edge (u, v) in G
 - Set these literals in S to be true
 - Then, the truth assignment will be consistent and all clauses are satisfied
- \Leftarrow :
 - Given a satisfying assignment, at least one literal from each triangle must be true so that every clause is satisfied
 - Pick one literal from each triangle to form an independent set S of size k



NP-completeness

Definitions

- P:** Decision problem X with poly-time algorithms
- NP:** Decision problems X with a poly-time certifier
 - Certifier $C(s, t)$ returns 'yes' for some certificate $|t| \leq p(|s|)$
- NP-complete:** $X \in NP$ and $\forall Y \in NP, Y \leq_p X$. Problem X is NP-complete if X is at least as hard as every NP problem

Properties

Proof: $P \subseteq NP$

- Given $X \in P$, there exists a poly-time algorithm A that returns $A(s)$.
- Construct a certifier $C(s, t)$ that returns $A(s)$ and set $t = \emptyset$
- Then, $C(s, t)$ runs in poly-time and $|t| \leq p(|s|)$

Proof: NP-complete problem X is solvable in polynomial time iff $P = NP$

\Rightarrow : X is in NP by definition. X is solvable in poly-time. Hence, $NP \subseteq P$

\Leftarrow : $P = NP$. X is in NP by definition. Hence, X is solvable in poly-time

Proof: If $\exists X \in NP$ s.t. X cannot be solved in polynomial time, then no NP-complete problem can be solved in polynomial time.

- Contrapositive of the above proof: If $P \neq NP$, then no NP-complete problem is solvable in polynomial time

Polynomial Transformations

- Cook reduction:** Problem X polynomial reduces to problem Y if X can be solved using (i) polynomial # of standard computational steps and (ii) polynomial # of calls to oracle that solves Y
- Karp reduction:** Problem X polynomial transforms to problem Y if given any input x to X, we can construct an input y s.t. x is a 'yes' instance of X iff y is a 'yes' instance of Y

Proof X is NP-complete

- If Y is NP-complete, $X \in NP$ and $Y \leq_p X$, then X is NP-complete:
 - Take an $Z \in NP$, then $Z \leq_p Y \leq_p X$. Hence, $Z \leq_p X$.
 - By defn of NP-completeness, X is NP-complete

Hence, it suffices to show:

Show that $X \in NP$

Choose a known NP-complete problem Y

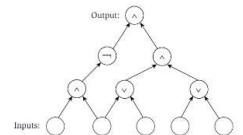
Prove $Y \leq_p X$: (Karp reduction)

Construct an arbitrary instance s_Y of Y and construct instance s_X of X in polynomial time

Show that $s_Y = \text{'yes'}$ iff $s_X = \text{'yes'}$

Circuit SAT \leq_p 3-SAT

- Circuit SAT:** Known NP-complete problem
- Reduce any $X \in NP$ to circuit SAT:
 - Set first n source nodes to be the n-bit s and remaining $p(|s|)$ source nodes to represent the bits in t
 - X outputs 'yes' iff the $\exists t$ s.t. $|t| \leq p(|s|)$ and the circuit is satisfiable (i.e. $C(s, t) = \text{'yes'}$)
- Show 3-SAT is NP-complete by reducing Circuit SAT to 3-SAT



Construction (Circuit SAT \leq_p 3-SAT)

- For each node v in Circuit SAT, let x_v be a variable in 3-SAT
- Note that $P \Rightarrow Q \Rightarrow \neg P \vee Q$
- NOT: $x_v \Leftrightarrow \bar{x}_u = (x_v \Rightarrow \bar{x}_u) \wedge (\bar{x}_u \Rightarrow x_v) = (\bar{x}_v \vee \bar{x}_u) \wedge (x_v \vee x_u)$
- AND: $x_v \Leftrightarrow (x_u \wedge x_w) = \dots = (\bar{x}_v \vee x_u) \wedge (\bar{x}_v \vee x_w) \wedge (x_v \vee \bar{x}_u \vee \bar{x}_w)$
- OR: $x_v \Leftrightarrow (x_u \vee x_w) = \dots = (x_v \vee \bar{x}_u) \wedge (x_v \vee \bar{x}_w) \wedge (\bar{x}_v \vee x_u \vee x_w)$
- SOURCE: Assign it to 0 or 1 if it is a hardcoded input
- OUTPUT: Add variable $x_o = 1$ for output
- Some clauses have < 3 variables. Create 4 new variables z_1, z_2, z_3, z_4 and add the clauses $(\bar{z}_1 \vee z_3 \vee z_4), (\bar{z}_1 \vee \bar{z}_3 \vee z_4), (\bar{z}_1 \vee z_3 \vee \bar{z}_4), (\bar{z}_1 \vee \bar{z}_3 \vee \bar{z}_4)$ for each $i = 1, i = 2$. This ensures that $z_1 = z_2 = 0$
- If a clause has 1 variable t , replace it with $(t \vee z_1 \vee z_2)$
- If a clause has 2 variables $(s \vee t)$, replace it with $(s \vee t \vee z_1)$

Proof

\Rightarrow Suppose Circuit SAT is satisfiable

- The satisfying assignment to Circuit SAT will create values at all nodes of the circuit
 - This set of values will satisfy the constructed SAT instance
- \Leftarrow Suppose 3-SAT is satisfiable
- The clauses in 3-SAT ensure that the values assigned to all nodes of the circuit are the same as what the circuit computes for these nodes.
 - $x_o = 1$ in 3-SAT, so the assignment is satisfiable in Circuit SAT

3-SAT \leq_p Hamiltonian Cycle

Construction

- Assume a 3-SAT instance Φ with n variables and k clauses,
- For each variable x_i in 3-SAT, construct a path $P_i = v_{i,1}, v_{i,2}, \dots, v_{i,b}$ where $b = 3k + 3$. Edges in the path are bi-directional
- For each path,
 1. Add edges from $v_{i,1}$ to $v_{i+1,1}$ and $v_{i+1,b}$
 2. Add edges from $v_{i,b}$ to $v_{i+1,1}$ and $v_{i+1,b}$
- Add s and edges from s to $v_{1,1}$ and $v_{1,b}$
- Add t and edges from $v_{n,1}$ and $v_{n,b}$ to t
- This models a Hamiltonian Cycle with 2^n paths (each 'layer' can be traversed from L2R or R2L, independent of other 'layers'). Similarly, there are 2^n different assignments in a 3-SAT instance
- If a path P_i is traversed from L2R, set $x_i = 1$. Else $x_i = 0$
- Add an extra node c_j for each clause C_j
 1. For each variable x_i in C_j , add the edges $x_{i,3j} \rightarrow c_j$ and $c_j \rightarrow x_{i,3j+1}$ if x_i is not a negation. Otherwise, add the edges $x_{i,3j+1} \rightarrow c_j$ and $c_j \rightarrow x_{i,3j}$
 2. e.g. if $C_1 = x_1 \vee \bar{x}_2 \vee x_3$, P_1 must go from L2R OR P_2 must go from R2L OR P_3 must go from L2R in order for c_j to be visited

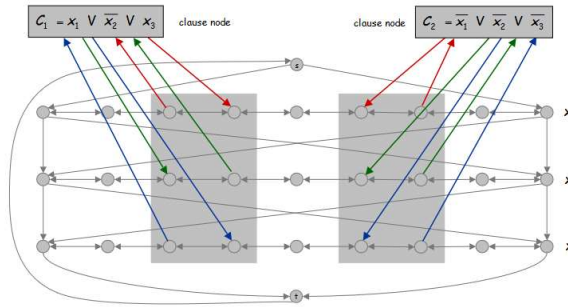
Proof

\Rightarrow Suppose Φ is satisfiable

- If an arbitrary variable x_i is 1, then P_i traverses from L2R, otherwise it traverses from R2L.
- For each clause C_j , since it is satisfied, there must be a path P_i that traverses in the "correct" direction so c_j can be spliced into the cycle via edges incident on $v_{i,3j}$ and $v_{i,3j+1}$

\Leftarrow Suppose there exists a Hamiltonian Cycle \mathcal{C}

- Then all c_j must be visited
- If \mathcal{C} enters a node c_j from $v_{i,3j}$, it must immediately depart on an edge to $v_{i,3j+1}$ otherwise $v_{i,3j+2}$ will never be visited without breaking the Hamiltonian property
- Symmetrically, if \mathcal{C} enters a node c_j from $v_{i,3j+1}$, it must immediately depart on an edge to $v_{i,3j}$ otherwise $v_{i,3j-1}$ will never be visited without breaking the Hamiltonian property
- Hence, for each c_j , the nodes before and after c_j in \mathcal{C} are joined by an edge in G .
- Therefore, we can remove each c_j in \mathcal{C} and join $v_{i,3j}$ and $v_{i,3j+1}$ via an edge, forming a Hamiltonian cycle \mathcal{C}' of $G' = G - \{c_1, \dots, c_k\}$
- Any Hamiltonian cycle in G' must traverse each P_i in only one direction. Hence, each P_i in \mathcal{C}' must traverse fully in only one direction
- If \mathcal{C}' traverses P_i from L2R, we can set $x_i = 1$, otherwise $x_i = 0$
- Since the larger cycle \mathcal{C} was able to visit all c_j , at least one of the paths was traversed in the correct direction relative to c_j
- Hence, the truth assignment satisfies all the clauses.



Hamiltonian Cycle \leq_p Traveling Salesman Problem

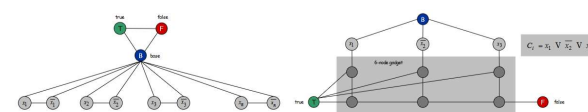
Construction

- Given graph $G(V, E)$, create $|V|$ cities with distance function
$$d(u, v) \begin{cases} 1, & (u, v) \in E \\ 2, & (u, v) \notin E \end{cases}$$
- Then, G will have a TSP tour of length $\leq |V|$ iff G is Hamiltonian

3-SAT \leq_p 3-Colorability

Construction

- For each variable x_i , create nodes x_i and \bar{x}_i and join them via an edge
- Create 3 extra nodes, T, F and B that connect to one another. Connect B to every literal
- For each clause, create a 6-node gadget:



- Φ is satisfiable iff the constructed graph is 3-colorable

Proof

\Rightarrow Suppose Φ is satisfiable

- Color all true literals green
 - Color nodes below green nodes to be red and the node below blue
 - Color remaining middle row nodes blue (first layer of nodes in gadget)
 - Color remaining bottom nodes red/green as forced
 - The resulting graph is 3-colorable
- \Leftarrow Suppose the graph is 3-colorable
- Assign each literal coloured green to be true
 - Each variable must have one literal to be green and the other to be red (refer to left image). Hence, the assignment is consistent
 - No literals can be blue, hence each literal is assigned T or F
 - At least one literal in any clause will be true (green) to satisfy 3-colorability (refer to right graph)
 - Hence, Φ is satisfiable

3-SAT \leq_p Subset-Sum

Construction

- Given Φ with n variables and k clauses, create $2n + 2k$ integers, each with $n + k$ digits
- For each integer, the first n digits represent each variable and the last k digits represent each clause
- For each variable x_i , create 2 integers for x_i and \bar{x}_i . For both integers, set digit i to 1. For each clause digit, set it to 1 if the literal is part of the clause.
- For each clause c_j , create $2k$ integers. Set first n digits to be 0. Set digit j to 1 for one of the integers and 2 for the other.
- Then, Φ is satisfiable iff there exists a subset that sums to $W = 11\dots144\dots4$ (n 1s and k 4s)

Proof

\Rightarrow Suppose Φ is satisfiable

- For each variable x_i , choose one of the two integers representing it (depending on where x_i is T or F)
- There will be n of such integers.
- For each clause c_j , choose one or both of the two integers representing it such that the digit j of the resulting sum is 4
- Sum the $n + k$ digits to get S
- The first n digits of S will be 1 because exactly one integer is picked for each integer
- The last k digits of S will be 4; because Φ is satisfiable, digit j has value 1, 2 or 3 from the sum of the n integers in (2). We can always pick some 1 or 2 of the dummy rows to make it up to 4
- Take the selected integers. Then, there exists a subset sum equals to W

\Leftarrow Suppose there exists a subset sum equals to W

- Let the subset of integers that sum to W be S
- Then, exactly one of the integers in S must have digit $i \leq n$ set to 1. To achieve this, exactly one of the literals for each variable in Φ must be true (therefore, there is consistent assignment)
- Each clause c_j is true in order for each digit $j \geq k$ to be 4
- Hence, Φ is satisfied

1. Let (A^*, B^*) be the optimal global min-cut of G and F^* be the set of edges across the cut. Let $|F^*| = k$
2. The first iteration of the contraction algorithm picks and contracts an edge in F^* with probability $\frac{k}{|E|}$
3. Suppose $\exists v \in V$ s.t. $\text{degree}(v) < k$. Take v to form A^* and the rest of the vertices to form B^* . Then, the global min-cut $< k$, a contradiction. Hence, $\forall v \in V, \text{degree}(v) \geq k$.

- For any graph, $|E| = \frac{1}{2} \sum_{v_i \in V} \text{degree}(v_i)$. Therefore, $|E| \geq \frac{1}{2} kn$
- Hence, by point (2), the algorithm contracts an edge in F^* in the first iteration with probability $\leq \frac{k}{|E|} \leq \frac{k}{\frac{1}{2} kn} = \frac{2}{n}$
- After j iterations, the number of vertices remaining is $n' = n - j$ (each iteration reduces the # vertices by one).
- Suppose no edge in F^* were contracted in these j iterations. Then, the min-cut is still k . Hence, $|E'| \geq \frac{1}{2} kn'$ and the algorithm contracts an edge in F^* with probability $\leq \frac{2}{n'}$
- Let E_j = event that an edge in F^* is not contracted in iteration j
- Then,

$$Pr(E_1 \cap E_2 \cap \dots \cap E_{n-2}) = Pr(E_1) \times Pr(E_2|E_1) \times Pr(E_3|E_1 \cap E_2) \times \dots \times Pr(E_{n-2}|E_1 \cap E_2 \cap \dots \cap E_{n-3}) \geq \left(1 - \frac{2}{n}\right) \left(1 - \frac{2}{n-1}\right) \left(1 - \frac{2}{n-2}\right) \dots \left(1 - \frac{2}{4}\right) \left(1 - \frac{2}{3}\right) = \frac{(n-2)}{n} \frac{(n-3)}{(n-1)} \frac{(n-4)}{(n-2)} \dots \frac{2}{4} \frac{1}{3} = \frac{2}{n(n-1)} \geq \frac{2}{n^2}$$
- Run the algorithm $n^2 \ln n$ times with independent random choices. Then, the probability of failing is

$$\leq \left(1 - \frac{2}{n^2}\right)^{n^2 \ln n} = \left(\left(1 - \frac{2}{n^2}\right)^{\frac{1}{2n^2}}\right)^{2 \ln n} \leq e^{-2 \ln n} = \frac{1}{n^2}$$

$$1 + x \leq e^x \Rightarrow 1 - 1/x \leq e^{-1/x} \Rightarrow (1 - 1/x)^x \leq e^{-1}$$

Expectation

- $E[X] = \sum_{j=0}^{\infty} j \Pr(X = j)$
- $X = \# \text{ trials until first success} = X \sim G(p) \Rightarrow E[X] = 1/p$
- $X = \text{Bernoulli Trial} = X \sim \text{Bernoulli}(p) \Rightarrow E[X] = p$
- Linearity of expectation (LoE): $E[X + Y] = E[X] + E[Y]$

Max 3-SAT (Las Vegas)

- Given Φ with k clauses, the expected # of clauses satisfied by a random assignment is $7k/8$
- Hence, there must exist a random assignment that satisfies at least $7k/8$ clauses with probability $\geq 1/8k$
- Johnson's algorithm: If we repeatedly generate random assignments until one satisfies $\geq 7k/8$ clauses, the number of trials is $\leq 8k$

Proof

- Let X_j = RV that equals 1 if clause j is satisfied and 0 otherwise
- Hence, $E[X_j] = p = 7/8$ (Bernoulli Trial)
- Let $X = \# \text{ of satisfied clauses i.e. } X \sim B(k, 7/8)$. Then, $E[X] = E[X_1 + X_2 + \dots + X_k] = E[X_1] + E[X_2] + \dots + E[X_k] = 7k/8$ (LoE)
- Let $p_j = \Pr[\# \text{satisfied} = j]$ and $p = \Pr[\# \text{satisfied} \geq 7k/8]$

Hence,

$$7k/8 = E[X] = \sum_{j \geq 0} j p_j = \sum_{j < 7k/8} j p_j + \sum_{j \geq 7k/8} j p_j$$

$$\leq \left(\frac{7k}{8} - \frac{1}{8}\right) \sum_{j < 7k/8} p_j + k \sum_{j \geq 7k/8} p_j$$

$$\leq \left(\frac{7k}{8} - \frac{1}{8}\right) (1) + kp$$

- Therefore, $p \geq 1/8k$
- Let $Y = \# \text{ trials until one assignment satisfies } \geq 7k/8 \text{ clauses. Then, } Y \sim G(p \geq 1/8k)$
- Hence, $E[Y] \leq 1/(1/8k) = 8k$

Universal Hashing

- Universal class of hash functions** = a set H of hash functions h_i s.t.
 $\Pr[h(u) = h(v)] \leq 1/n, \quad \forall u, v \in U, n = \# \text{ buckets}$
- Let $X = \# \text{ collisions with any } u \in U \text{ using } H$. Then, for any $S \subseteq U$ and $|S| \leq 1/n, E_{h \in H}[X] \leq 1$

Proof

- Let $X_v = \text{RV that equals 1 if } v \text{ collides with } u \text{ and 0 otherwise}$
- Hence, $E[X_v] = \Pr[h(u) = h(v)] \leq 1/n$
- Let $X = \# \text{ collisions with any } u \text{ i.e. } X \sim B(|S|, p \leq 1/n)$. Then, $E[X] \leq |S|/n \leq 1$

- Designing such a universal class:

- Choose a prime number $p, n \leq p \leq 2n$ (Chebyshev: such a p exists)
- For each $\forall u \in U$, identify a base- p integer of r digits: $x = (x_1, x_2, \dots, x_r)$
- Let A be the set of all possible r -digit, base- p integers i.e. for each $a \in A, a = (a_1, a_2, \dots, a_r), 0 \leq a_i < p$
- Then, $H = \{h_a \mid h_a(x) = (\sum_{i=1}^r a_i x_i) \bmod p, a \in A\}$

Proof

- Let $x = (x_1, x_2, \dots, x_r)$ and $y = (y_1, y_2, \dots, y_r), x \neq y$
- Then, $\exists j$ s.t. $x_j \neq y_j$
- $h_a(x) = h_a(y) \Leftrightarrow a_1(y_1 - x_1) + \dots + a_m(y_m - x_m) \bmod p$
- Assume vector a is chosen randomly by first (uniformly) randomly picking each $a_i, i \neq j$ in the range $[0, p]$, then picking a_j at random
- Since p is prime, $a_j z = m \bmod p$ has at most (exactly) 1 solution among p possibilities
 - Let p be prime, $z \neq 0 \bmod p$ (z not divisible by p).
 - Suppose α and β are 2 different solutions (for contradiction)
 - Then, $\alpha z = m + k_1 p$ and $\beta z = m + k_2 p$. So, $(\alpha - \beta)z = (k_1 - k_2)p = 0 \bmod p$
 - Hence, $(\alpha - \beta)$ is divisible by p since z not divisible by p
 - Therefore, $\alpha = \beta$ since $0 \leq \alpha, \beta < p$ (contradiction)
- Hence, $\Pr[h_a(x) = h_a(y)] \leq 1/p \leq 1/n$ since $n \leq p \leq 2n$

Chernoff Bounds

- Let X_1, X_2, \dots, X_n be a Bernoulli process (i.e. they are independent 0-1 RVs) and $X = X_1, X_2, \dots, X_n$. Then, $\forall \mu \geq E[X]$ and $\forall \delta > 0$,

$$\Pr[X > (1 + \delta)\mu] < \left(\frac{e^\delta}{(1 + \delta)^{1+\delta}}\right)^\mu$$

and $\forall \mu \leq E[X]$ and $\forall \delta \in (0, 1)$,

$$\Pr[X < (1 - \delta)\mu] < e^{-\delta^2 \mu / 2}$$

Proof (above mean)

- Markov's inequality states that $\Pr[X > a] \leq E[X]/a$
- Hence, $\Pr[X > (1 + \delta)\mu] = \Pr[e^{tX} > e^{t(1+\delta)\mu}] \leq e^{-t(1+\delta)\mu} E[e^{tX}]$
- $E[e^{tX}] = E[e^{t \sum_i X_i}] = E[e^{tX_1} e^{tX_2} \dots e^{tX_n}] = \prod_i E[e^{tX_i}]$ (independence: $X \perp Y \Rightarrow E(XY) = E(X)E(Y)$)
- Let $p_i = \Pr[X_i = 1]$. Then,
 $E[e^{tX_i}] = p_i e^t + (1 - p_i) e^0 = 1 + p_i(e^t - 1) \leq e^{p_i(e^t - 1)}$
- Therefore,
 $\Pr[X > (1 + \delta)\mu] \leq e^{-t(1+\delta)\mu} \prod_i E[e^{tX_i}] \leq e^{-t(1+\delta)\mu} \prod_i e^{p_i(e^t - 1)} \leq e^{-t(1+\delta)\mu} e^{\mu(e^t - 1)} \quad (\mu \geq E[X] = \sum_i p_i)$
- Take $t = \ln(1 + \delta)$. Then,

$$\Pr[X > (1 + \delta)\mu] \leq e^{-\ln(1+\delta)(1+\delta)\mu} e^{\mu(e^{\ln(1+\delta)} - 1)} = \left(\frac{e^\delta}{(1 + \delta)^{1+\delta}}\right)^\mu$$

Load Balancing

- m jobs arrive in a stream and need to be processed immediately by n identical processors. Assign jobs to processors uniformly at random without centralised controller (without round-robin, where each processor will receive at most $\lceil m/n \rceil$ jobs. How likely is it that some processor is assigned "too many" jobs?

Analysis

- Let $X_i = \# \text{ jobs assigned to processor } i$
- Let $Y_{ij} = \text{RV that equals 1 if job } j \text{ is assigned to processor } i$
- Then, $E[Y_{ij}] = 1/n$ ($E(X) = p$, if $X \sim \text{Bernoulli}(p)$)
- Hence, $X_i = \sum_j Y_{ij} \Rightarrow \mu = E[X_i] = nE[Y_{ij}] = 1$
- By Chernoff bounds, with $\delta = c - 1$, $\Pr[X_i > c] < \frac{e^{c-1}}{c^c}$ for some c which will be chosen later
- Let $\gamma(n) = x$ s.t. $x^x = n$ and choose $c = e\gamma(n)$
- Then, $\Pr[X_i > c] < \frac{e^{c-1}}{c^c} < \left(\frac{e}{c}\right)^c = \left(\frac{1}{\gamma(n)}\right)^{e\gamma(n)} < \left(\frac{1}{\gamma(n)}\right)^{2\gamma(n)} = \frac{1}{n^2}$
- By Union Bound, $\Pr[\cup_{i=1}^n X_i > c] \leq n \left(\frac{1}{n^2}\right) = \frac{1}{n}$. That is, the probability that at least one processor received more than $c = e\gamma(n)$ jobs is $\frac{1}{n}$.
- Therefore, the probability that no processors received more than $c = e\gamma(n) = \Theta(\log n / \log \log n)$ jobs is $1 - \frac{1}{n}$
- Suppose there are $m = 16n \ln n$ jobs. Then, $E[X_i] = 16 \ln n$. With high probability, every processor will have between half and twice the average load (e.g. $8 \ln n = \frac{1}{2} \mu \leq X_i \leq 2\mu = 16 \ln n$)

Proof

- By Chernoff bounds, with $\delta = 1$,

$$\Pr[X_i > 2\mu] < \left(\frac{e}{4}\right)^{16 \ln n} < \left(\frac{1}{e}\right)^{2 \ln n} = \frac{1}{n^2}$$

$$\Pr\left[X_i < \frac{1}{2}\mu\right] < e^{-\frac{1}{2}\left(\frac{1}{2}\right)^2 (16 \ln n)} = \frac{1}{n^2}$$

- By Union Bound, $\Pr[\cup_{i=1}^n X_i > 2\mu] \leq n \left(\frac{1}{n^2}\right) = \frac{1}{n}$ and $\Pr[\cup_{i=1}^n X_i < \frac{1}{2}\mu] \leq n \left(\frac{1}{n^2}\right) = \frac{1}{n}$.
- Therefore, the probability that every processor has load between half and twice the average load is $\geq 1 - \frac{1}{n} - \frac{1}{n} = 1 - \frac{2}{n}$